



# Network Socket Programming - 2

---

BUPT/QMUL

2010-10-19



# Review

---

- Basic Concepts in NP
  - Introduction to Network Programming
    - Importance
    - Classes
    - Environments in this course
  - Program Developing
    - Phases
    - Skills
    - Useful tools
  - Basic Concepts
    - Process
    - File Descriptor
    - System Call
    - Signal



# About the labs on 2010-10-15

---


- Open the file with `O_APPEND`
  - Read OK, Write failed.
- Open the file with `O_RDWR|O_APPEND`
  - Read OK
  - Write at the end of file




# Agenda

---

- Basic concepts in NP
- Introduction to IP & TCP/UDP
- *Introduction to Sockets*



# Introduction to IP & TCP/UDP




# Introduction to Sockets



# Introduction to Sockets

---

- Reviews of some helpful points
- Sockets interface
- Major system calls
- Sample programs



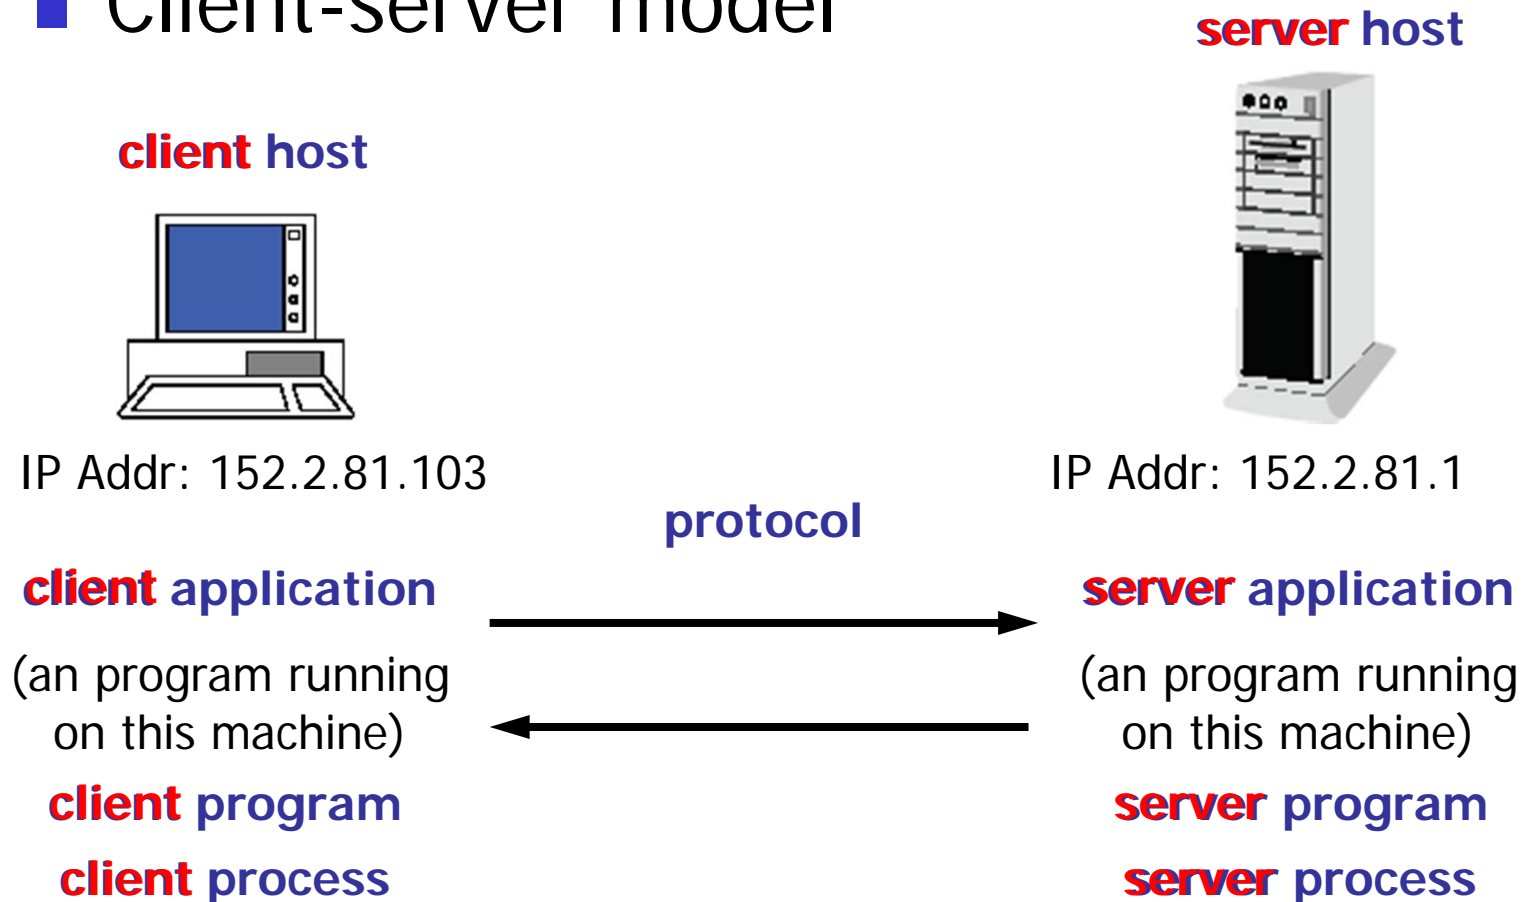
# Introduction to Sockets

## Part I: some helpful points



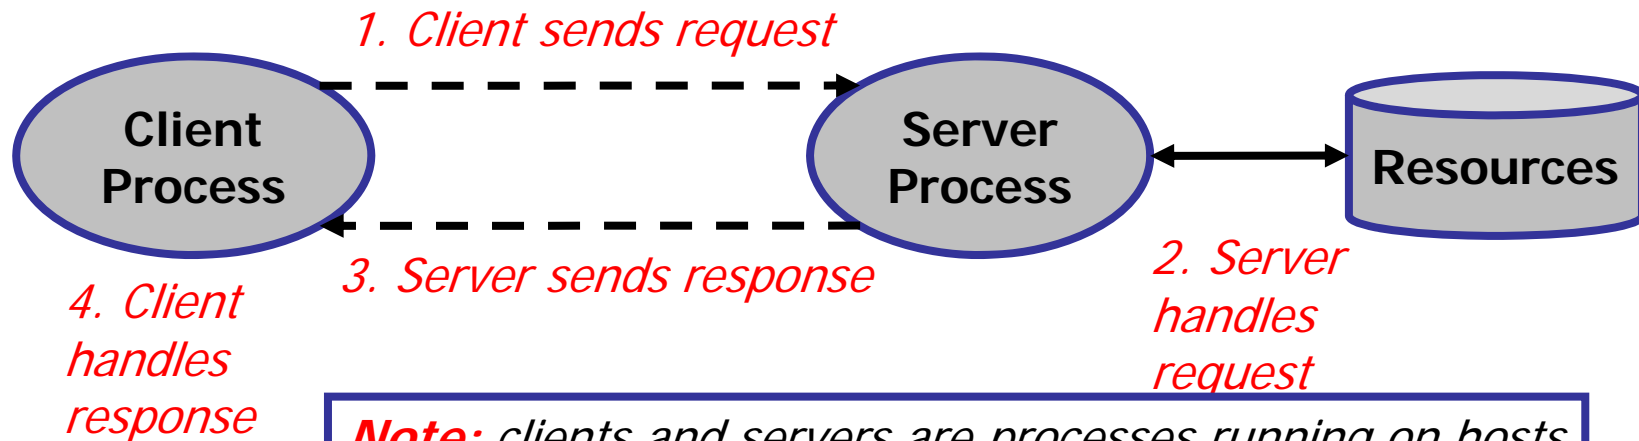
# Reviews Of Some Helpful Points

## ■ Client-server model



# A Client-server Transaction


- Most of network applications are based on the client-server model:
  - A **server** process and one or more **client** processes
  - Server manages some **resources**.
  - Server provides **service** by manipulating resources for clients.



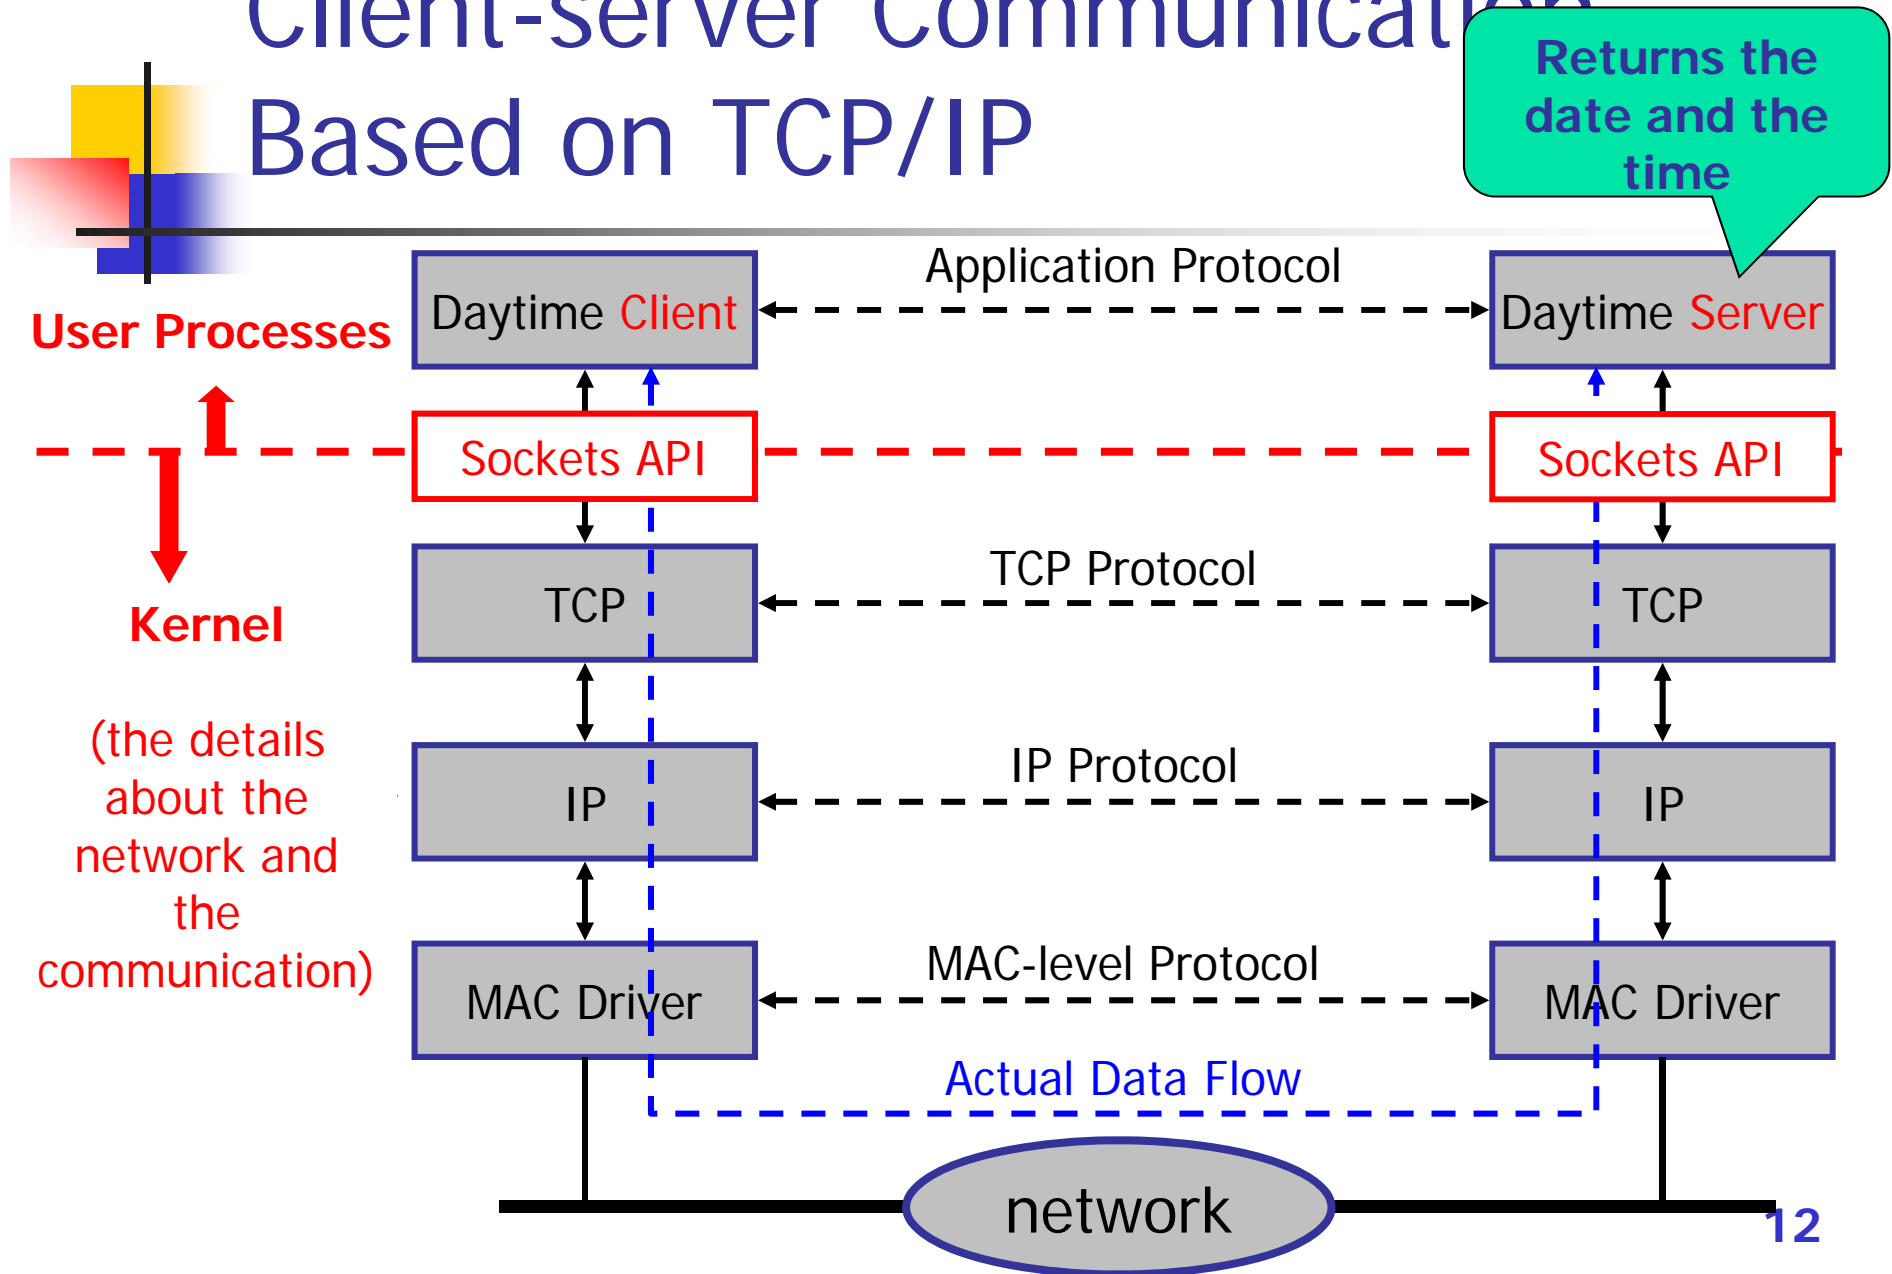
**Note:** clients and servers are processes running on hosts (can be the same or different hosts).

# Reviews Of Some Helpful Points

## ■ Communication on the Internet



# Client-server Communication Based on TCP/IP





# Reviews Of Some Helpful Points

---

- A programmer's view of the Internet
  - 1. Hosts are mapped to a set of 32-bit *IP addresses*.
    - 202.112.96.163
  - 2. The set of IP addresses is mapped to a set of identifiers called Internet *domain names*.
    - 202.112.96.163 is mapped to www.bupt.edu.cn
  - 3. A process on one Internet host can communicate with a process on another Internet host over a *connection*.

# IP Addresses (1)

- 32-bit IP addresses are stored in an *IP Address structure*


- <netinet/in.h>

```
/* Internet address. */  
typedef uint32_t in_addr_t;  
struct in_addr {  
    in_addr_t s_addr;  
};
```

```
/*Defined in <stdint.h>*/  
typedef unsigned int uint32_t;
```

- Two ways to store multi-byte integers

- Big-endian vs. little-endian





## IP Addresses (2)

---

- Host byte order vs. network byte order
  - Host byte order is machine-dependent
    - You can see it in `<bits/endian.h>`
    - A program used to output the host byte order
  - Network byte order is machine-independent (*big-endian*)
  - Byte order conversion functions
    - `htonl`: convert **long int** from host to network byte order.
    - `htons`: convert **short int** from host to network byte order.
    - `ntohl`: convert **long int** from network to host byte order.
    - `ntohs`: convert **short int** from network to host byte order

# IP Addresses (3)

-- A program used to output the host byte order

## ■ byteorder.c

```
#include <sys/types.h>
#include <stdio.h>

int main(int argc, char **argv)
{
    union {
        short s;
        char c[sizeof(short)];
    } un;

    un.s=0x0102;

    if (sizeof(short)==2){
        if (un.c[0]==1 &&un.c[1]==2)
            printf("big-endian\n");
        else if (un.c[0]==2 &&un.c[1]==1)
            printf("little-endian\n");
        else
            printf("Unknow\n");
    }else
        printf("sizeof(short)=%d\n",sizeof(short));

    exit(0);
}
```

a **union** is like a **struct**, only all the data members sit at the same memory location. This means only one of them can be used at a time.



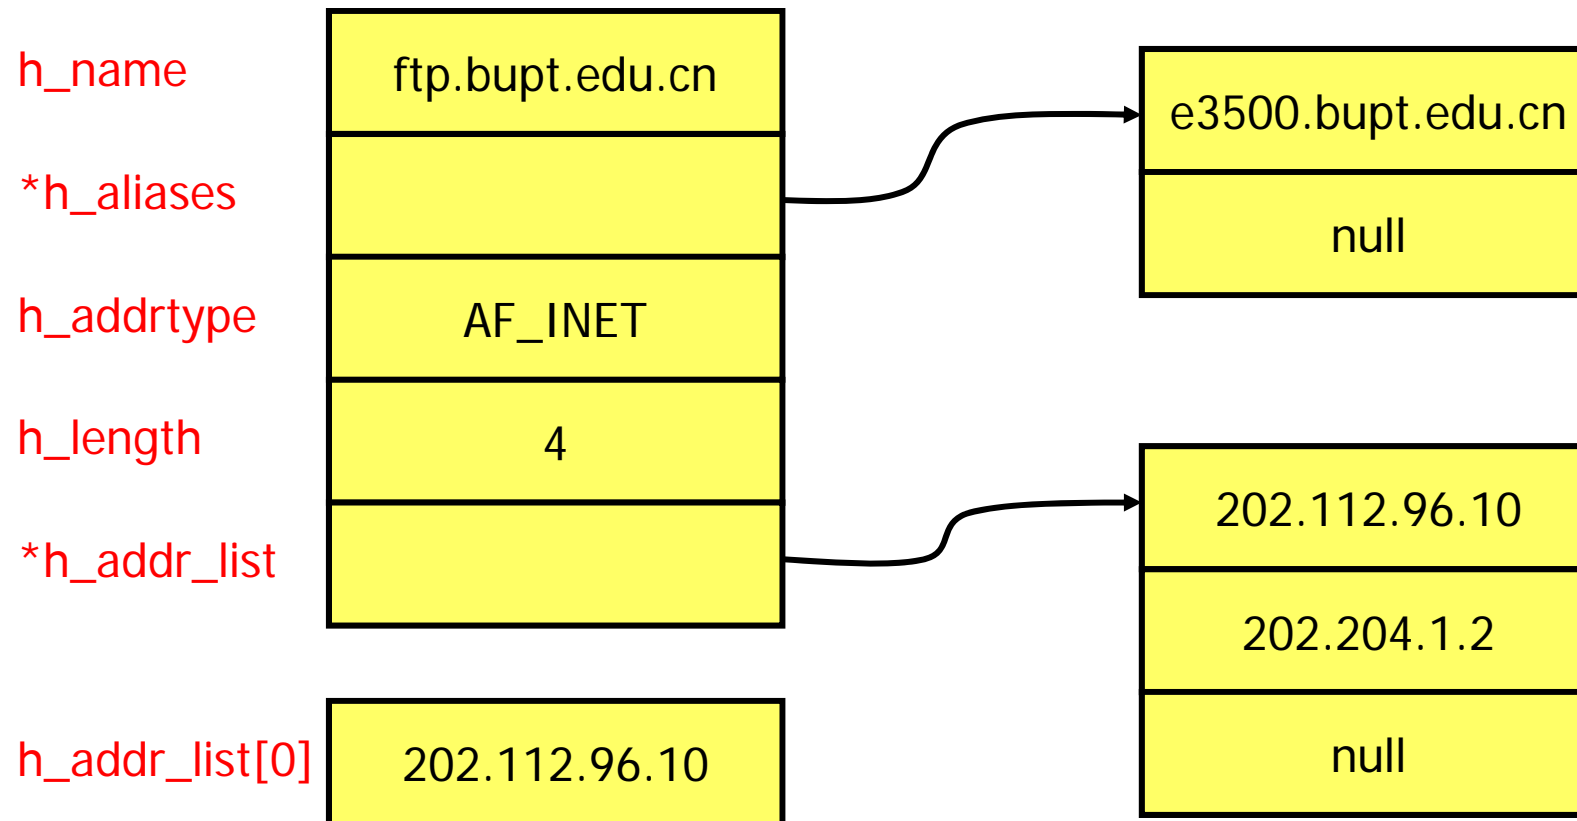


# Domain Name System (1)

- The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called *DNS*.
  - Conceptually, programmers can view the DNS database as a collection of millions of *host entry structures*
  - <netdb.h>

```
/* Description of data base entry for a single host. */
struct hostent
{
    char *h_name;           /* Official name of host. */
    char **h_aliases;       /* Alias list. */
    int h_addrtype;         /* Host address type. */
    int h_length;           /* Length of address. */
    char **h_addr_list;     /* List of addresses from name server. */
#define h_addr h_addr_list[0] /* The first address in the address list. */
};
```

## Domain Name System (2)





## Domain Name System (3)

- Functions for retrieving host entries from DNS
  - gethostbyname: query key is a DNS domain name.

```
#include <netdb.h>
struct hostent * gethostbyname (const char *hostname);
```

- gethostbyaddr: query key is an IP address.

```
#include <netdb.h>
struct hostent * gethostbyaddr (const char *addr, int len, int family);
```

↓  
**AF\_INET**  
**for IPv4**

# How to print IP address?

- Conversion between numeric address and text string

```
#include <sys/socket.h>
```

```
#include <arpa/inet.h>
```

```
const char *inet_ntop(int af, const void *src, char *dst, socklen_t cnt)
```

```
int inet_pton(int af, const char *src, void *dst)
```

↓  
**AF\_INET**  
for IPv4

↓  
**string**

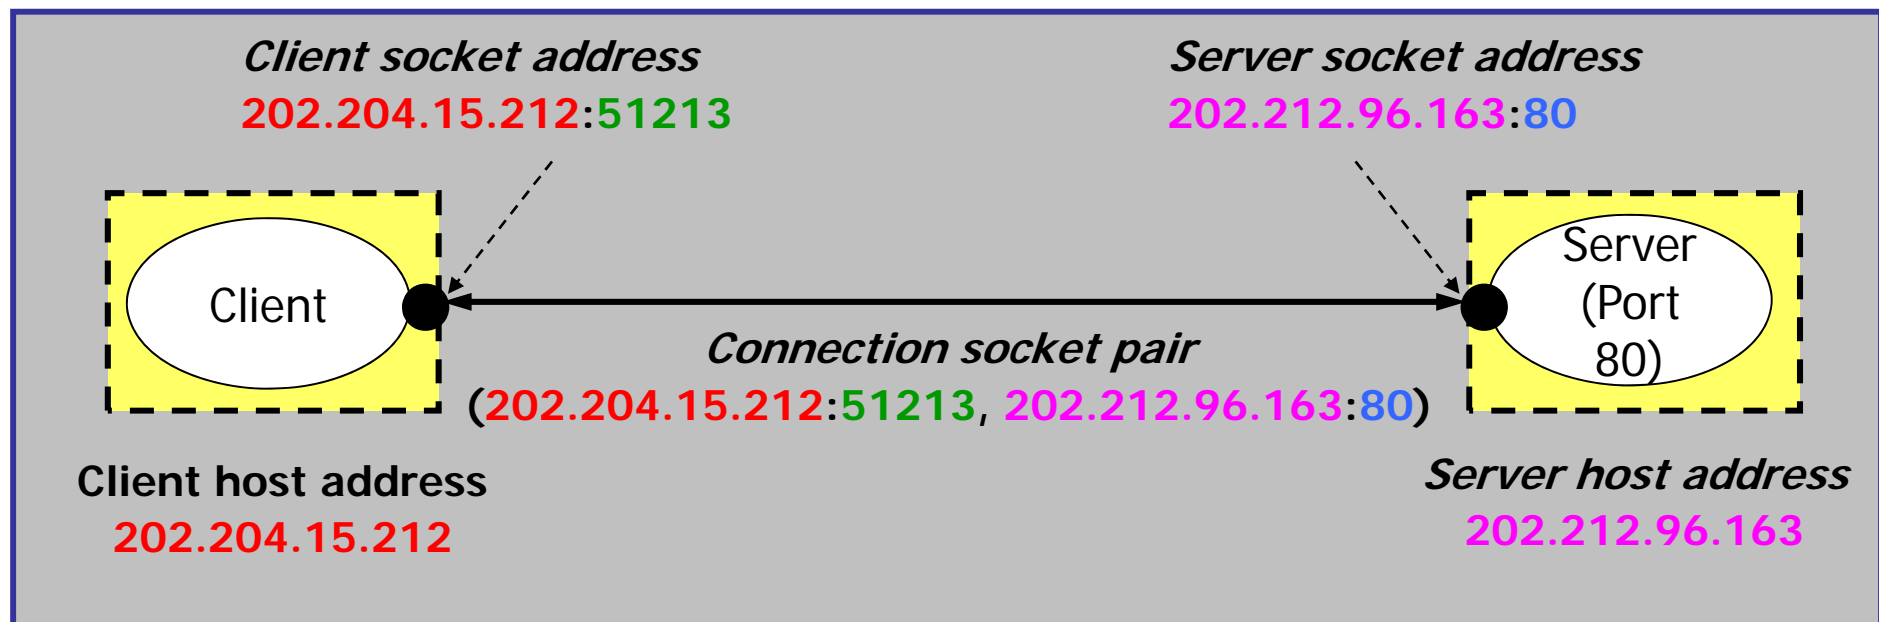
↓  
**network address**

↓  
**string**

↓  
**length of address**

# Connections

- Clients and servers communicate by sending streams of bytes over *connections*.
- Connections are end-to-end, full-duplex (2- way communication), and reliable.





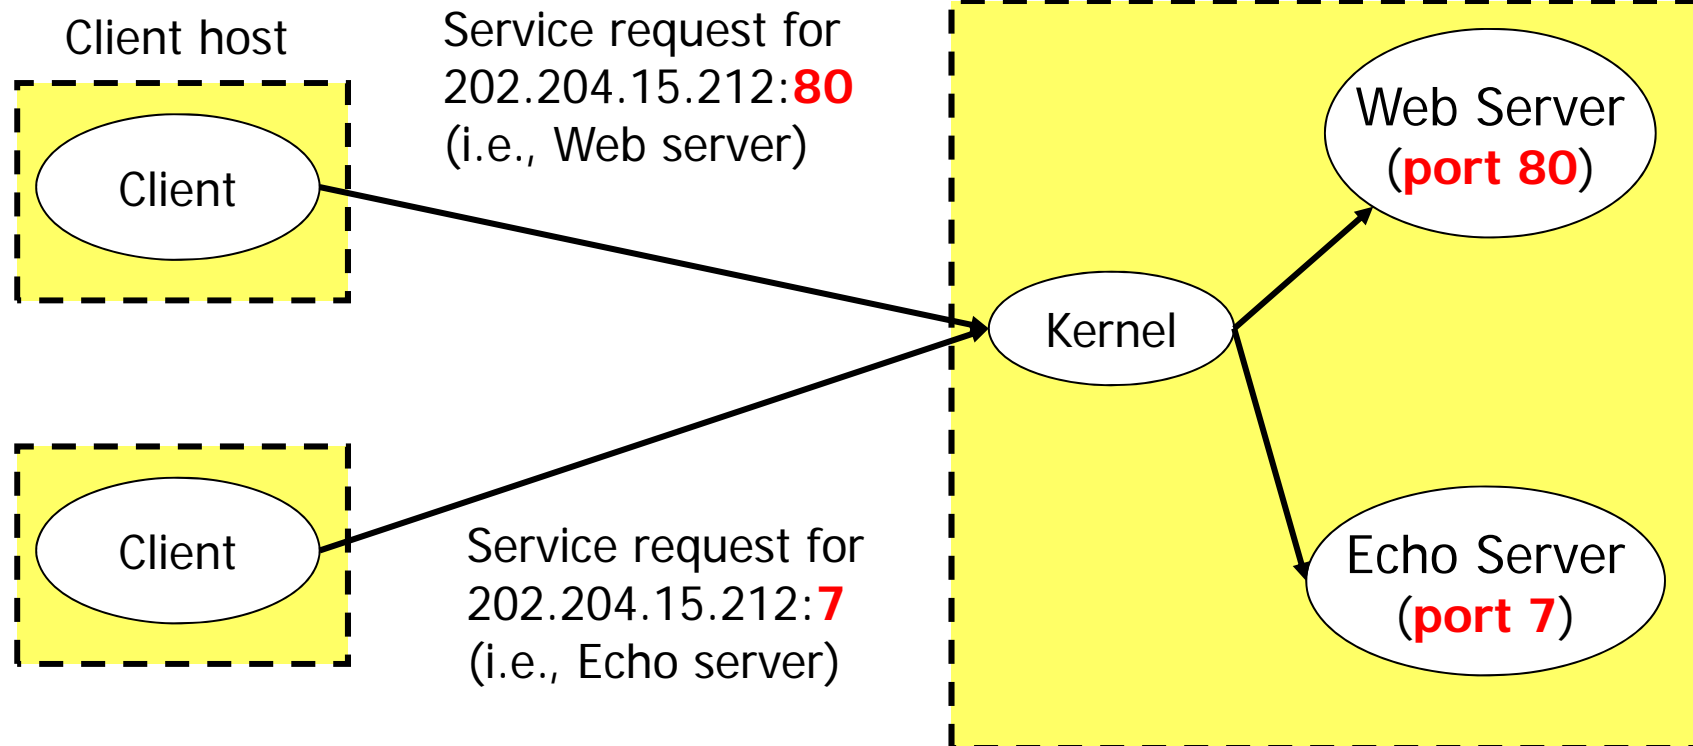
# Clients

---

- Examples of client programs
  - Web browsers, ftp, telnet, ssh
- How does a client find the server?
  - The IP address in the server socket address identifies the host (*more precisely, an adapter on the host*)
  - The **(well-known) port** in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
  - Examples of well know ports
    - Port 7: Echo server
    - Port 23: Telnet server
    - Port 25: Mail server
    - Port 80: Web server

# Using ports to identify services

Server host 202.204.15.212





# Servers

---

- Servers are long-running processes (daemons).
  - Typically created at boot-time by the **init process** (process 1#)
  - Run continuously until the machine is turned off
- Each server waits for requests to arrive on a well-known port associated with a particular service.
  - See ***/etc/services*** for a comprehensive list of the services available on a Linux machine
- A machine that runs a server process is also often referred to as a “server.”





# Server examples

Name	Port	Services	Resources
Web server	80	Retrieves files and runs CGI programs on behalf of the client	files/compute cycles (CGI programs)
FTP server	20, 21	stores and retrieve files	files
TELNET server	23	proxies a terminal on the server machine	terminal
Mail server	25	stores mail messages in spool file	email “spool” file



# Useful Unix Commands

---

- netstat
- ifconfig
- ping



# Useful Unix Commands - netstat

---

- Functions: prints information about the Linux networking subsystem, e.g., network connections, routing tables, interface statistics etc.
- **netstat**
  - Displays a list of open sockets.
- **netstat -i**
  - Display the information about the network interfaces
- **netstat -ni**
  - Display the information about the network interfaces using numeric addresses
- **netstat -r**
  - Display the kernel routing tables
- **netstat -nr**
  - Display the kernel routing tables using numeric addresses

# Useful Unix Commands - netstat

## ■ netstat

```
[root@localhost include]# netstat
Active Internet connections (w/o servers)

```

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State
tcp	0	0	192.168.1.253:telnet	192.168.1.27:3256	ESTABLISHED
tcp	0	0	::ffff:192.168.1.253:ssh	::ffff:192.168.1.27:2888	ESTABLISHED
tcp	0	0	::ffff:192.168.1.253:ssh	::ffff:192.168.1.27:3047	ESTABLISHED

```
Active UNIX domain sockets (w/o servers)

```

Proto	RefCnt	Flags	Type	State	I-Node	Path
unix	10	[ ]	DGRAM		5724	/dev/log
unix	2	[ ]	DGRAM		6859	@/var/run/hal/hotplug_socket
unix	2	[ ]	DGRAM		3351	@udev
unix	2	[ ]	DGRAM		927082	
unix	2	[ ]	DGRAM		926850	
unix	3	[ ]	STREAM	CONNECTED	924266	
unix	3	[ ]	STREAM	CONNECTED	924265	
unix	3	[ ]	STREAM	CONNECTED	916866	/tmp/.X11-unix/X16



# Useful Unix Commands - netstat

## ■ netstat -ni

```
[root@localhost ~]# netstat -ni
Kernel Interface table
```

Iface	MTU	Met	RX-OK	RX-ERR	RX-DRP	RX-OVR	TX-OK	TX-ERR	TX-DRP	TX-OVR	Flg
eth0	1500	0	49922499	0	0	0	20779220	0	0	0	BMRU
lo	16436	0	1934	0	0	0	1934	0	0	0	LRU

- Ethernet interface is called **eth0** or **le0** depending on the machine
- Loop back interface is called **lo** and the common used IP address is 127.0.0.1



# Useful Unix Commands - netstat

## ■ netstat -nr

```
[root@localhost ~]# netstat -nr
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	MSS	Window	irtt	Iface
192.168.1.0	0.0.0.0	255.255.255.0	U	0	0	0	eth0
169.254.0.0	0.0.0.0	255.255.0.0	U	0	0	0	eth0
0.0.0.0	192.168.1.1	0.0.0.0	UG	0	0	0	eth0



# Useful Unix Commands - ifconfig

- Functions: configure the network interfaces, and usually be used to print the configuration of the network interfaces

```
[root@localhost ~]# ifconfig
eth0 Link encap:Ethernet HWaddr 00:13:72:4F:9D:3A
      inet addr:192.168.1.253 Bcast:192.168.1.255 Mask:255.255.255.0
      inet6 addr: fe80::213:72ff:fe4f:9d3a/64 Scope:Link
      UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
      RX packets:49923781 errors:0 dropped:0 overruns:0 frame:0
      TX packets:20779648 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:1000
      RX bytes:647355456 (617.3 MiB) TX bytes:2713364 (2.5 MiB)
      Base address:0xecc0 Memory:fe6e0000-fe700000

lo Link encap:Local Loopback
      inet addr:127.0.0.1 Mask:255.0.0.0
      inet6 addr: ::1/128 Scope:Host
      UP LOOPBACK RUNNING MTU:16436 Metric:1
      RX packets:1934 errors:0 dropped:0 overruns:0 frame:0
      TX packets:1934 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:0
      RX bytes:266858 (260.6 KiB) TX bytes:266858 (260.6 KiB)
```



# Useful Unix Commands - ping

- Functions: Sends a packet to the host specified by destination and prints out the roundtrip time ( Uses ICMP messages)


```
[root@localhost etc]# ping 192.168.1.27
PING 192.168.1.27 (192.168.1.27) 56(84) bytes of data.
64 bytes from 192.168.1.27: icmp_seq=0 ttl=128 time=0.261 ms
64 bytes from 192.168.1.27: icmp_seq=1 ttl=128 time=0.219 ms
64 bytes from 192.168.1.27: icmp_seq=2 ttl=128 time=0.181 ms

--- 192.168.1.27 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 1999ms
rtt min/avg/max/mdev = 0.181/0.220/0.261/0.034 ms, pipe 2
```

```
[root@localhost etc]# ping www.baidu.com
PING www.a.shifen.com (202.108.22.5) 56(84) bytes of data.
64 bytes from xd-22-5-a8.bta.net.cn (202.108.22.5): icmp_seq=0 ttl=57 time=363 ms
64 bytes from xd-22-5-a8.bta.net.cn (202.108.22.5): icmp_seq=1 ttl=57 time=177 ms
64 bytes from xd-22-5-a8.bta.net.cn (202.108.22.5): icmp_seq=2 ttl=57 time=172 ms

--- www.a.shifen.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 1999ms
rtt min/avg/max/mdev = 172.446/237.748/363.698/89.081 ms, pipe 2
```





# Introduction to Sockets

## Part II: sockets interface



# Sockets Interface

---

- Functions
- Definitions
- Types



# Sockets Interface – functions

---

- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols
- Provides a user-level interface to the network
- Underlying basis for all Internet applications
- Based on client/server programming model



# Sockets Interface – definitions(1)

---

- What is a socket?
  - To the kernel, a socket is an **endpoint** of communication.
  - To an application, a socket is a **file descriptor** that lets the application read/write from/to the network.
    - **Remember**: All Unix I/O devices, including networks, are modeled as files.
- Clients and servers communicate with each other by **reading from and writing** to **socket descriptors**.
- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors.



# File Model in Unix/Linux

---

- In Unix/Linux, all I/O devices are treated as files
  - Identified with File Descriptors
  - File operations
    - open
    - close
    - lseek
    - read
    - write
    - ...

File Descriptor Table  
(One per Process)

0	stdin
1	stdout
2	stderr
3	file
4	device
5	socket
..	...



# Sockets Interface – definitions(2)

- **Internet-specific socket address** (bits/socket.h)

```
struct sockaddr_in {  
    unsigned short sin_family; /* address family (always AF_INET) */  
    unsigned short sin_port; /* port num in network byte order */  
    struct in_addr sin_addr; /* IP addr in network byte order */  
    unsigned char sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```

- **Address family**: Domains refer to the area where the communicating processes exist. Commonly used domains include:
  - **AF\_UNIX**: for communication between processes on one system;
  - **AF\_INET** (IPv4): for communication between processes on the same or different systems using the DARPA standard protocols (IP/UDP/TCP)
  - **AF\_INET6** (IPv6)
  - **AF\_LOCAL** (Unix domain)
  - **AF\_UNSPEC** (the importance will be explained later)
  - ...



# Sockets Interface – definitions(3)

- **Generic socket address** (<sys/socket.h>)

```
struct sockaddr {  
    unsigned short sa_family; /* protocol family */  
    char sa_data[14]; /* protocol-specific address,  
                        up to 14 bytes. */  
};
```

- **Protocol family**
  - **PF\_LOCAL**: Local to host, pipes and file-domain
  - **PF\_UNIX**: Old BSD name for PF\_LOCAL
  - **PF\_INET**: **IP protocol family**
  - **PF\_AX25**: Amateur radio AX.25
  - **PF\_IPX**: Novell internet protocol
  - **PF\_INET6**: IP version 6
  - **PF\_ATMSVC**: ATM SVCs
  - **PF\_APPLETALK**: Appletalk DDP
  - ...



## Sockets Interface – definitions(4)

- Generic socket address and Internet-specific socket address
  - Pointer to generic socket address is used for address arguments to **connect, bind, and accept**
  - Must **cast** Internet-specific socket address (`sockaddr_in *`) to generic socket address (`sockaddr *`) for connect, bind, and accept

```
struct sockaddr_in serv;  
/* fill in serv{ } */  
bind (sockfd, (struct sockaddr *)&serv, sizeof(serv));
```





# Sockets Interface – types(1)

---

- Stream Socket
  - Service: reliable (i.e. sequenced, non-duplicated, non-corrupted) bidirectional delivery of byte-stream data
  - Metaphor: a phone call
  - `int s = socket (PF_INET, SOCK_STREAM, 0);`
- Datagram Socket
  - Service: unreliable, unsequenced datagram
  - Metaphor: sending a letter
  - `int s = socket (PF_INET, SOCK_DGRAM, 0);`
- Raw Sockets Service
  - allows user-defined protocols that interface with IP
  - Requires ``root`` access
  - Metaphor: playing with an erector set
  - `int s = socket (PF_INET, SOCK_RAW, protocol);`
- **SOCK\_STREAM and SOCK\_DGRAM are the most common types of sockets used within UNIX**



# Sockets Interface – types(2)

---


- Reliably-delivered Message Socket
  - Service: reliable datagram
  - Metaphor: sending a registered letter
  - Similar to datagram socket but ensure the arrival of the datagrams
  - `int s = socket (PF_NS, SOCK_RDM, 0);`
- Sequenced Packet Stream Socket
  - Service: reliable, bi-directional delivery of recordoriented data
  - Metaphor: record-oriented TCP
  - Similar to stream socket but using fixed-size datagrams
  - `int s = socket (PF_NS, SOCK_SEQPACKET, 0);`



# Abbreviations


---

<b>API</b>	Application Programming Interface
<b>CGI</b>	Common Gateway Interface
<b>DNS</b>	Domain Name System
<b>FTP</b>	File Transfer Protocol
<b>ICMP</b>	Internet Control Message Protocol
<b>IP</b>	Internet Protocol
<b>MAC</b>	Media Access Control
<b>TCP</b>	Transport Control Protocol
<b>UDP</b>	User Datagram Protocol



# Introduction to Sockets

## Part III: major system calls



# Introduction to Sockets

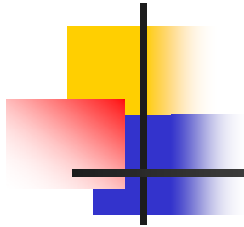
## Part IV: sample programs



## Labs on 2010-10-22

---

- Find the fore mentioned header files in your system
  - in.h, types.h, netdb.h, endian.h, socket.h, ...
- Find the host byte order of your machine
- Use man to learn the usage of
  - netstat, ifconfig, ping and traceroute



- Write a program to find the DNS information about a given host
  - The host may be specified in domain name or IP address, e.g.,
    - ./<exefile> www.baidu.com
    - ./<exefile> 202.108.22.5
  - Use gethostbyaddr() and gethostbyname()
  - Your program shall list the official name, all the aliases, all the IP addresses in numbers-and-dots format