



DNS Basics

BUPT/QMUL

2010-11-09



Related Information


- Basic function of DNS
- Host entry structure in Unix
- Two system calls for DNS database retrieving
 - `gethostbyname ()`
 - `gethostbyaddr ()`



Agenda

- Brief introduction to DNS
- Elements of the DNS
- DNS services
- DNS Protocols
- DNS tools

Refer to Chapter 23 of textbook



Brief Introduction to DNS

- Basic functions of DNS
- A short history of DNS




Basic Functions of DNS (1)

- Generally, applications refer to hosts/mailboxes and network resources by **ASCII strings** - such as
 - www.bupt.edu.cn
 - www.qmul.ac.uk
 - webmaster@company-a.com
- Nevertheless, the network itself only understands **binary addresses**, so some mechanism is required to convert the ASCII string to network addresses and vice versa.
- Low-level name: IP address
- High-level name: hostname




Basic Functions of DNS (2)

- Translating between addresses
 - **Hostname** (www.company-a.com)
 - Convert between Domain Name and IP
 - **IP address** (128.9.32.254)
 - Convert between IP and MAC
 - **MAC address** (C2-67-3E-95-C5-F1)



A Short History (1)

- Hostnames growth trend, by Nov. 2010



(From: <http://news.netcraft.com/>)



A Short History (2)

– from flat namespace to hierarchical namespace

- Original status

- there was simple a file, *hosts.txt*, that listed all the hosts and their IP addresses

flat structure

- Every night all the hosts would collect this file from the host that maintained it

Centralized control

- Incurred problems

- Each name had to be unique because the namespace was flat
- Excessive access to the machine maintaining the list
- Difficult to maintain when the network grew

Not scalable



A Short History (3)

– from flat namespace to hierarchical namespace


- Nowadays status

- *Hierarchical structure*
- *Distributed database*
- *Efficient, reliable, general purpose*
- It is primarily used for **mapping** host names and email destinations to IP addresses - but can be used for **other purposes**
- It is a **query / response** protocol running on top of **UDP/TCP**, with default port number **53**



Summary: What is DNS?

- Domain Name System
- A **distributed database** providing mapping between Domain name and IP address
 - Implemented in hierarchy of many **name servers**
- An **application protocol**
 - Used by hosts and name servers to communicate to **resolve names**



Elements Of The DNS



Elements Of DNS

- Domain namespace and resource records
- Name servers
- Name resolvers
- Protocol

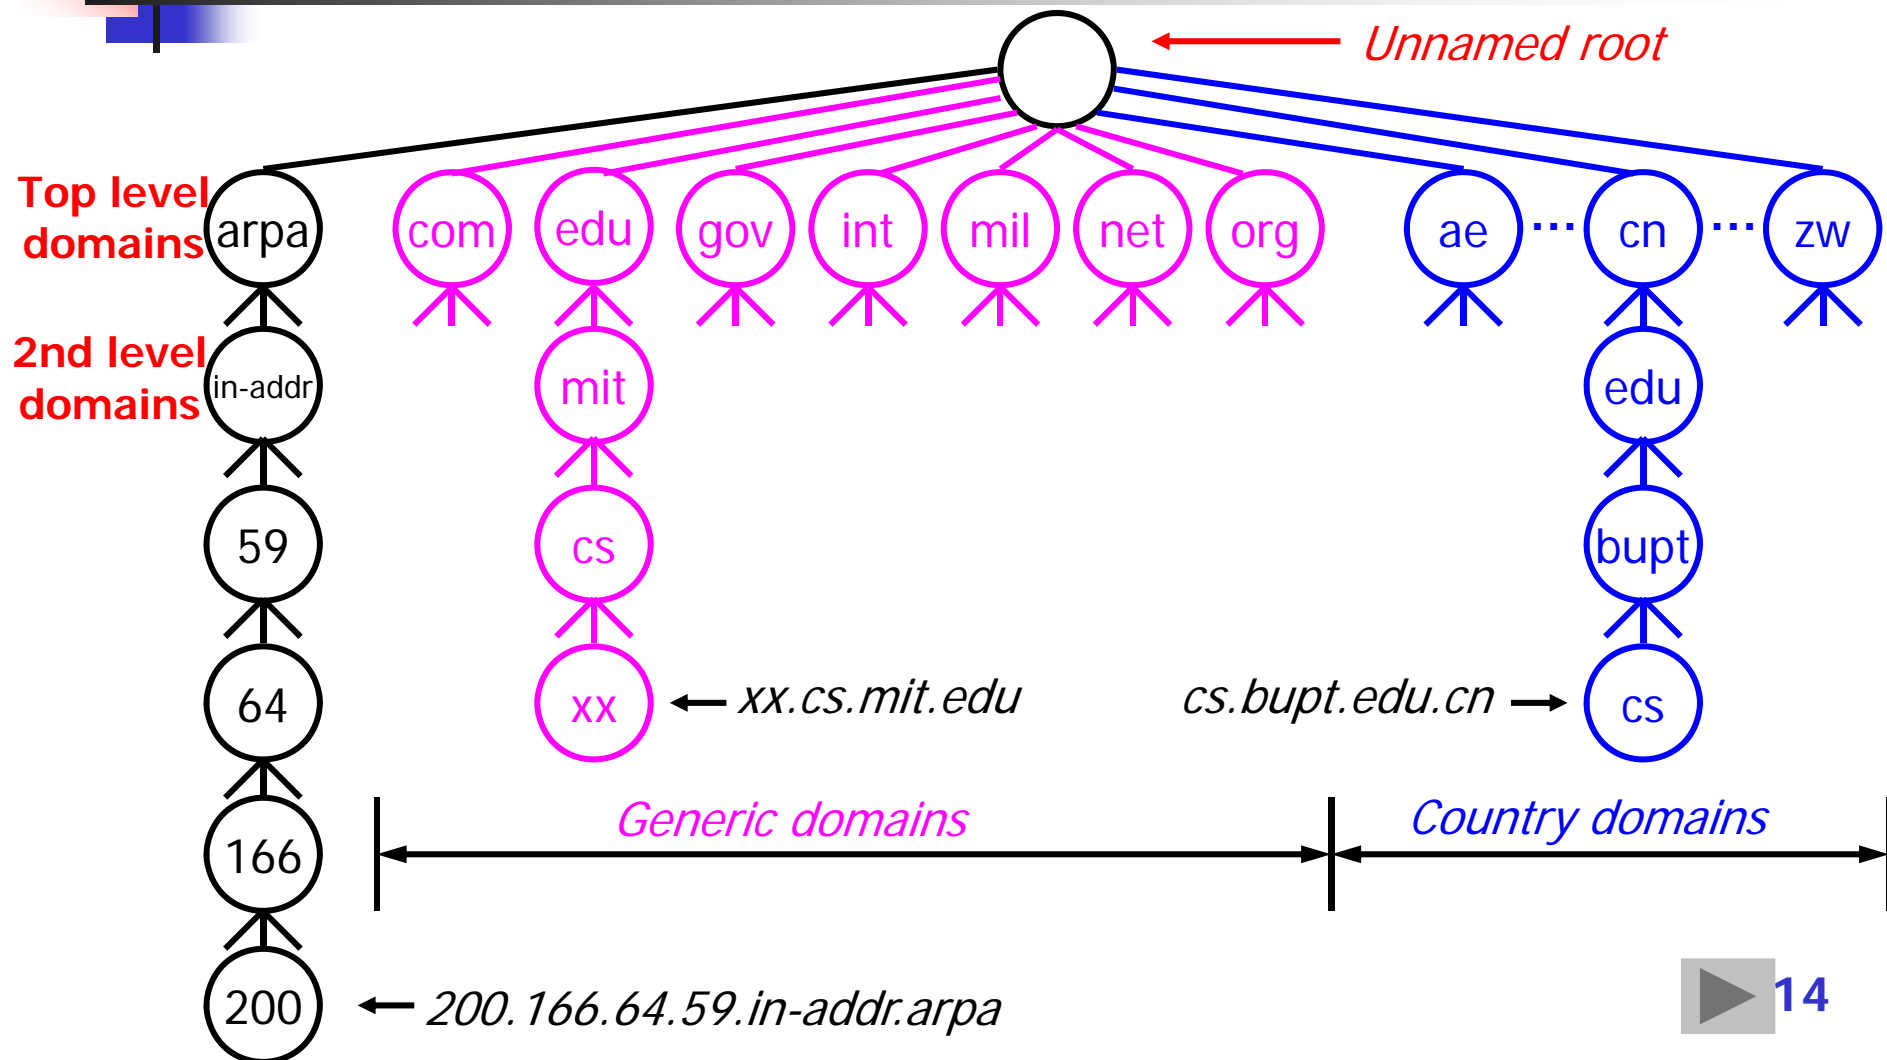


Domain namespace and resource records

- **Domain namespace** is of a hierarchical structure like an upended tree
- **RR** (Resource Record) is the data associated with a particular name

Domain Namespace (1)

– the hierarchical structure





Domain Namespace (2)

– organization of domains

- Each element of the hierarchy is referred to as a **domain**
- At the top of the hierarchy is the **root domain**, known as simply “.”
- Subdomains directly underneath the root domain are called **top-level domains**
- Domains directly underneath top-level domains are called **second-level domains**, and so on



Domain Namespace (3)

– the domain names

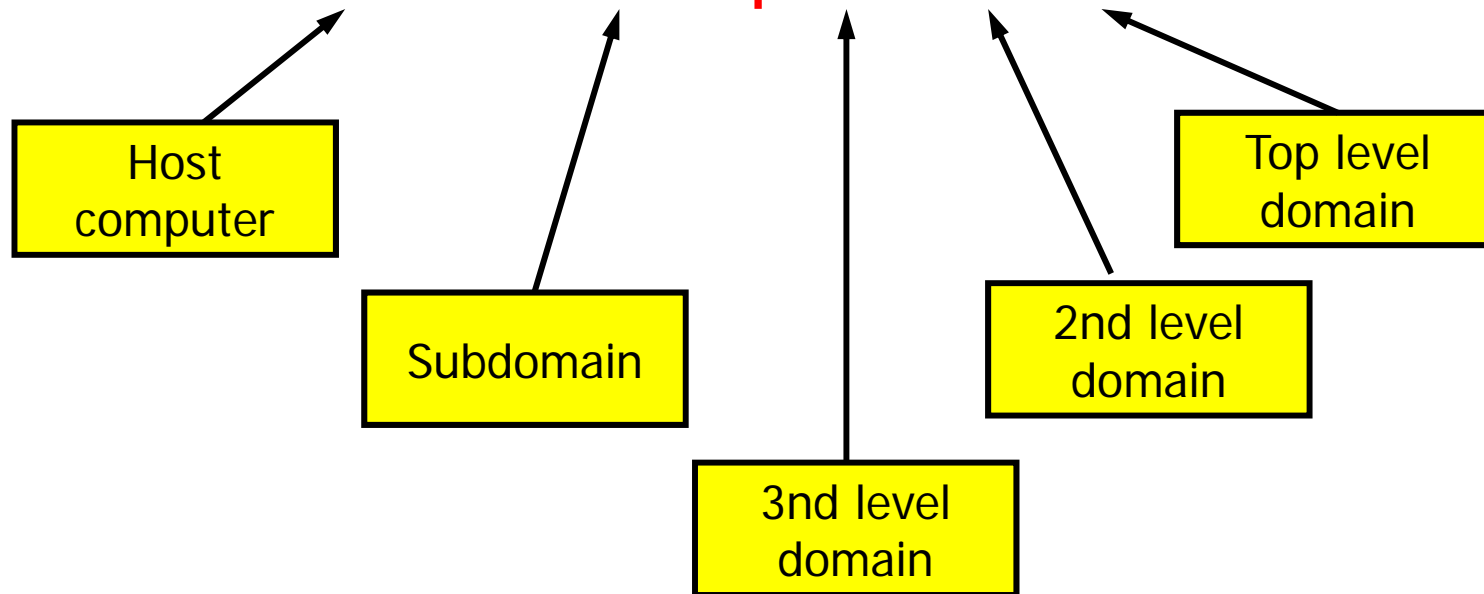
- The domain name of a node is **the list of the labels** on the path from the node to the root of the tree separated by dots(".")
- Conventionally printed or read **left to right**, from the most specific (lowest, farthest from the root) to the least specific (highest, closest to the root)
 - eg. www.bupt.edu.cn
- **Case insensitive**
- Components can be up to **63** characters long, the full pathname must not be more than **255** characters
- The full name of a domain is also called its **Fully Qualified Domain Name (FQDN)**

Domain Namespace (4)

– the domain names

- Example of domain name

chimera.elec.qmul.ac.uk





Resource Record

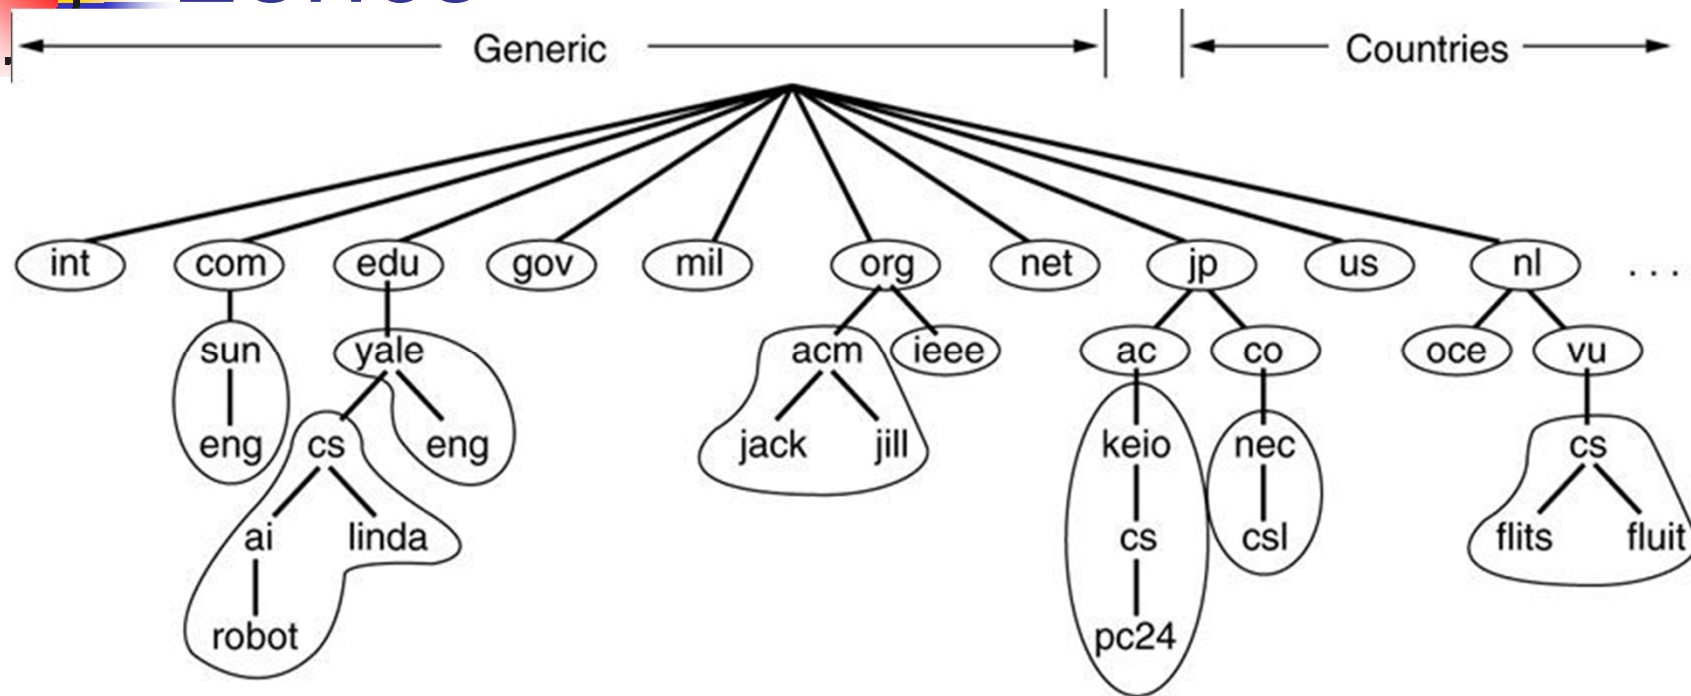
- Each domain in the DNS has one or more **Resource Records** (RRs), which are fields that contain information about that domain
- Each RR has the following information
 - **Owner**: the domain name where the RR is found
 - **Type**: specifies the type of the resource in this RR
 - A – Host Address
 - MX – Mail Exchanger
 - ...
 - **Class**: specifies the protocol family to use
 - IN – the Internet system
 - **TTL**: specifies the Time To Live (in unit of second) of the cached RRs
 - **RDATA**: the resource data



Name Servers (1)

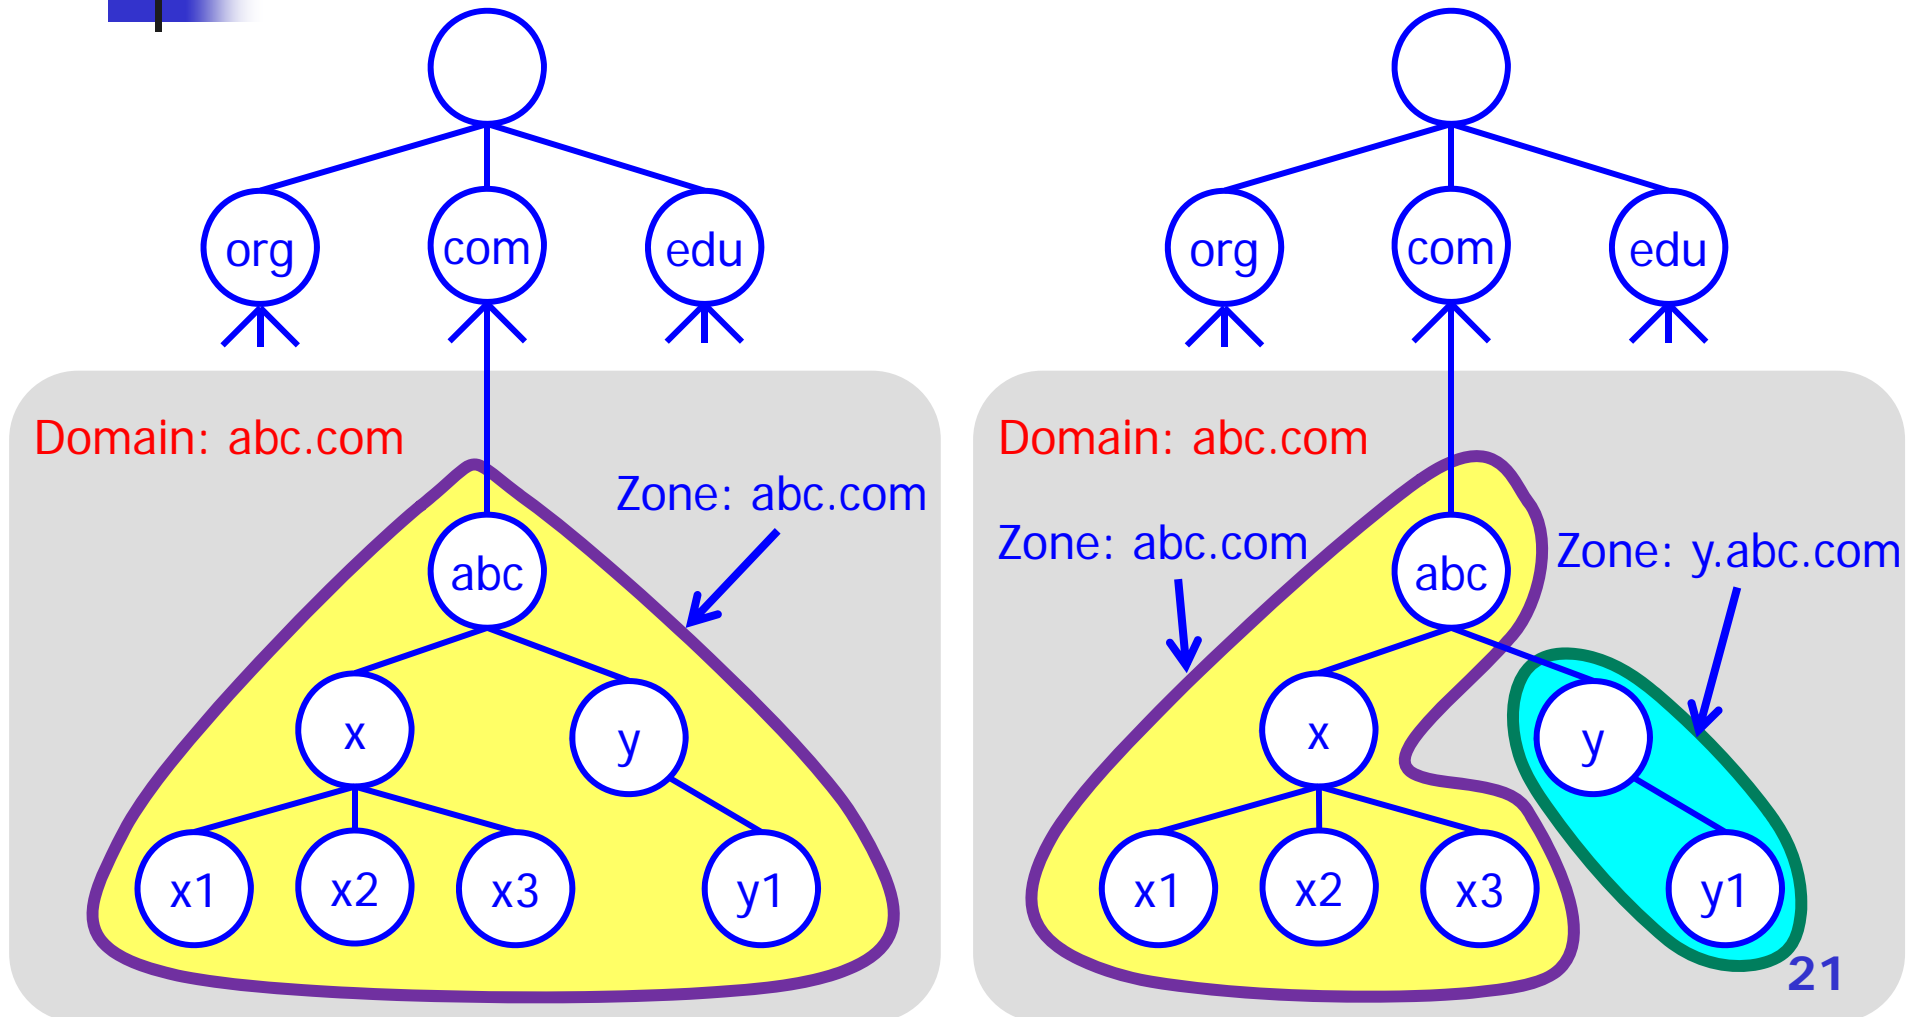
- Name servers are the **repositories** of information that make up the domain database.
- The database is divided up into sections called **zones**, which are distributed among the name servers. A zone may be one or more domains or even a sub-domain
- Each name server handles **one or more** zones. And the essential task of a name server is to **answer queries** using data in its zones.
- Name servers can answer queries in a simple manner. The response can always be generated **using only local data**, and either contains **the answer to the question** or **a referral to other name servers** "closer" to the desired information.
- A given zone will be **available from several name servers** to ensure its availability.

Zones



- A **zone** corresponds to an **administrative** authority that is responsible for that portion of the hierarchy
- Eg. **bupt** controls *x.bupt.edu.cn*

Example of Zone and domain





Name Servers (2)

- **Primary server / Authoritative server**
 - holds in its database the name-to-address mappings for the group of hosts it administers
 - knows the **official** answer
- **Secondary server**
 - maintains a copy of the Primary Server's database
- **Caching server**
 - asks DNS queries to other servers but maintains a cache of the responses together with a “time to live” value
 - **non-authoritative** data about other parts of the tree

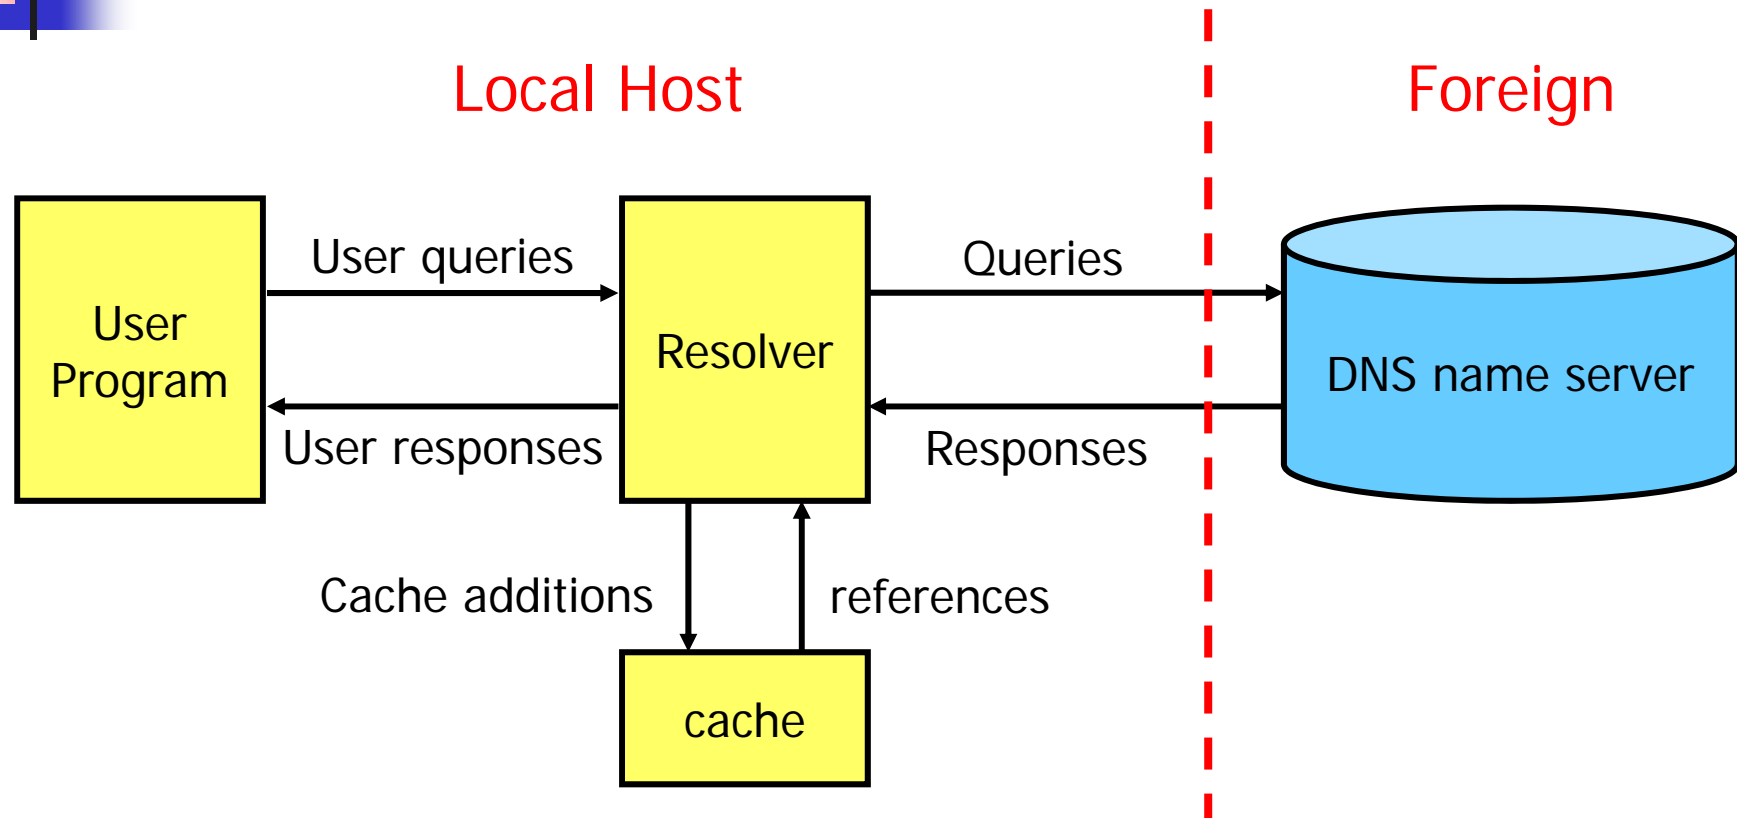


Name Resolvers (1)

- The **client** side of DNS
- A resolver is the **interface** between the user program and the domain name servers
 - A resolver receives a request from a user program (e.g., mail programs, TELNET, FTP)
 - asks questions to the DNS system on behalf of the application
 - returns the desired information

Name Resolvers (2)

– communication model

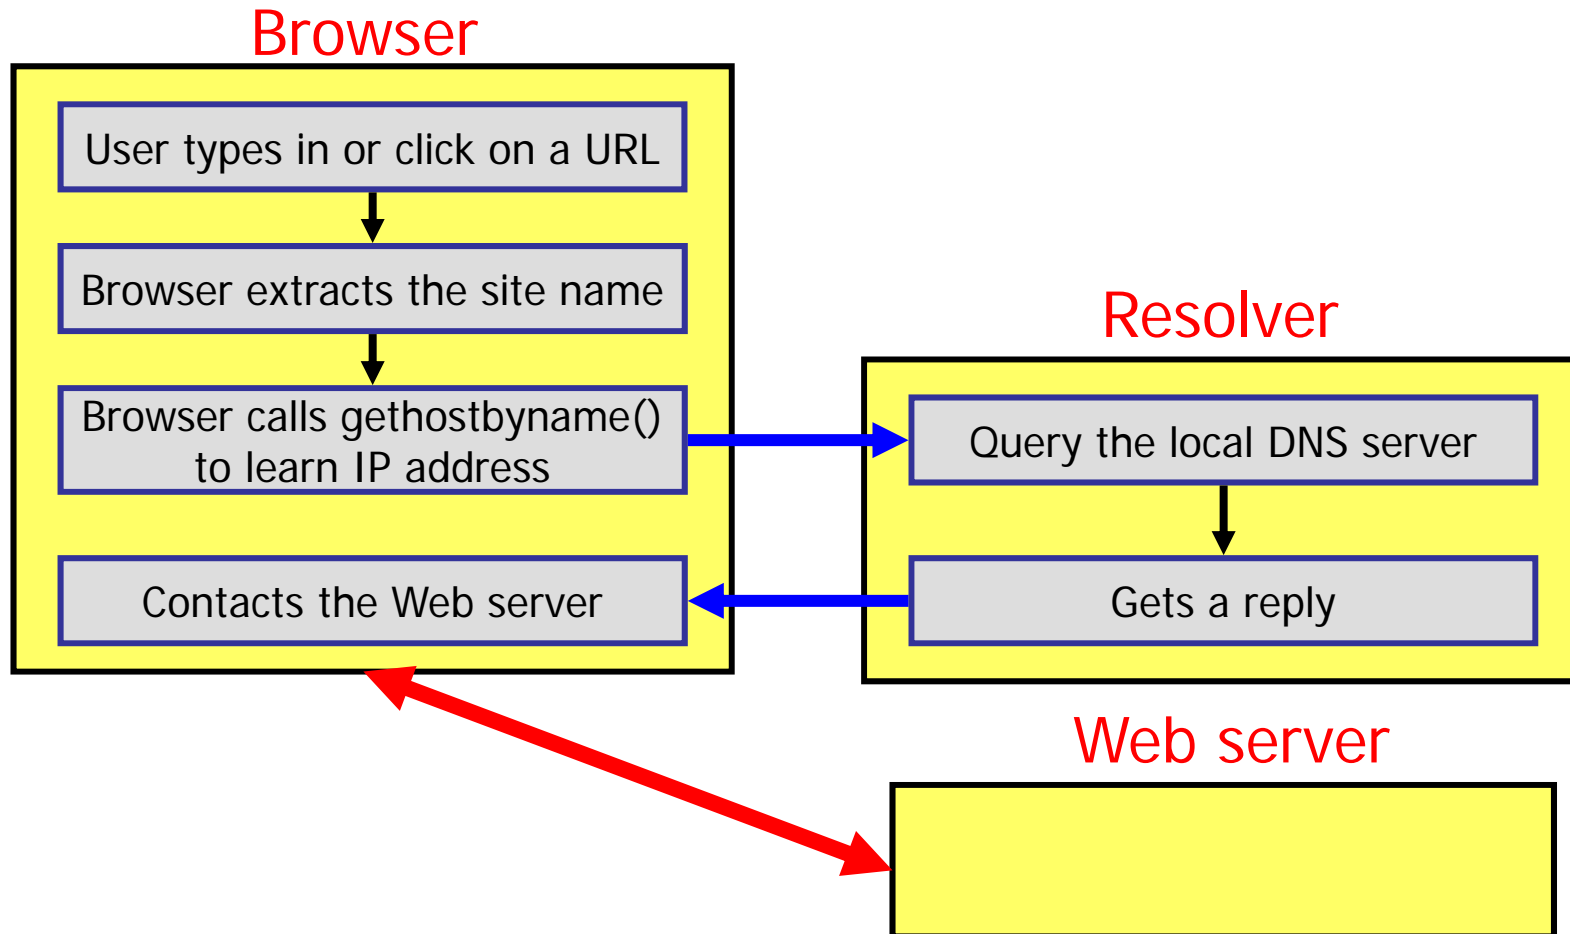



Resolvers are normally implemented in system calls.
E.g. `gethostbyname()`, `gethostbyaddr()`

Name Resolvers (3)

– an example

- DNS working together with HTTP application





DNS Services




DNS Services

- Mapping domain name to addresses
- Inverse queries (optional)
- Pointer queries

Mapping Domain Names to Addresses (1)

– arrangement of name servers

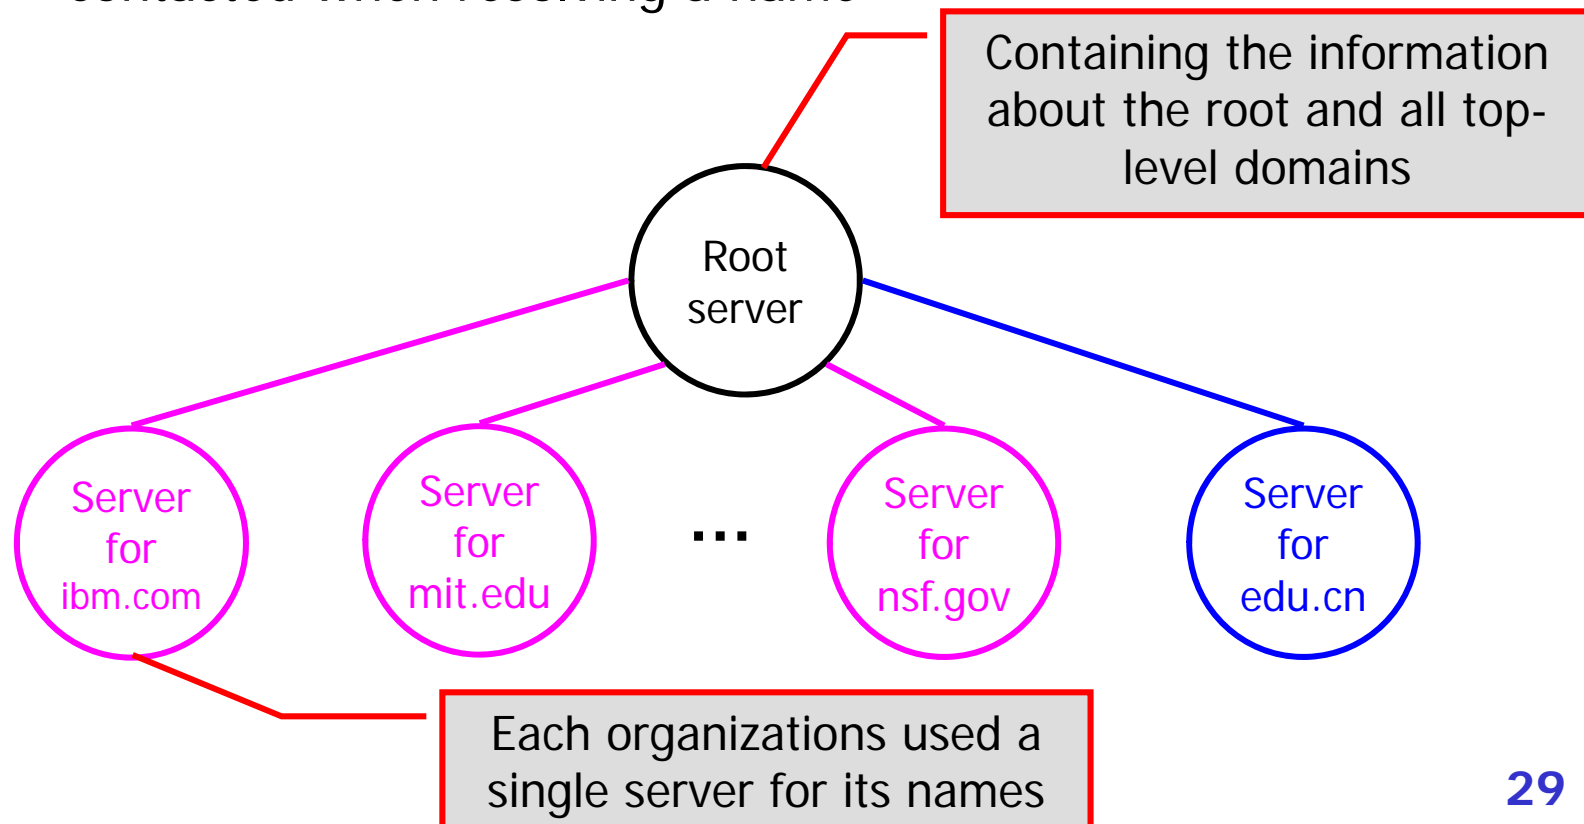
- Conceptual arrangement of name servers in a **tree** that corresponding to the name hierarchy
- Each server knows the addresses of all **lower-level** servers for all **subdomains** within the domain it handles



Mapping Domain Names to Addresses (2)

– arrangement of name servers

- A more realistic organization of name servers
- The tree is broad and flat and fewer servers need to be contacted when resolving a name





Mapping Domain Names to Addresses (3)

– name resolution process

- **Two-step** name resolution process
 - Beginning with the **local name server** (default name server)
 - If the local server can not resolve a name, the query must be sent to another server in the domain system
- It can improve the query efficiency because most queries to name servers refer to local name.




Mapping Domain Names to Addresses (4)

– two name resolution methods

- A query is made to a **local name server**. If local name server can not resolve the name, then it will query another server.
- **Recursive resolution**:
 - If the **queried server** does not have the information, it must **make** the appropriate **query** or queries to get the information
 - Generally, a server fulfills a recursive query either with data in its own memory or by making another recursive query
- **Iterative resolution**: If the **queried server** does not have the information, it may then **respond with** the **address of another server**; the **local name server** (on behalf of resolver)then **queries** that server (which might respond with the address of another server, and so on)
 - Commonly used by name servers on the **Internet**


Mapping Domain Names to Addresses (5)

– example of recursive resolution

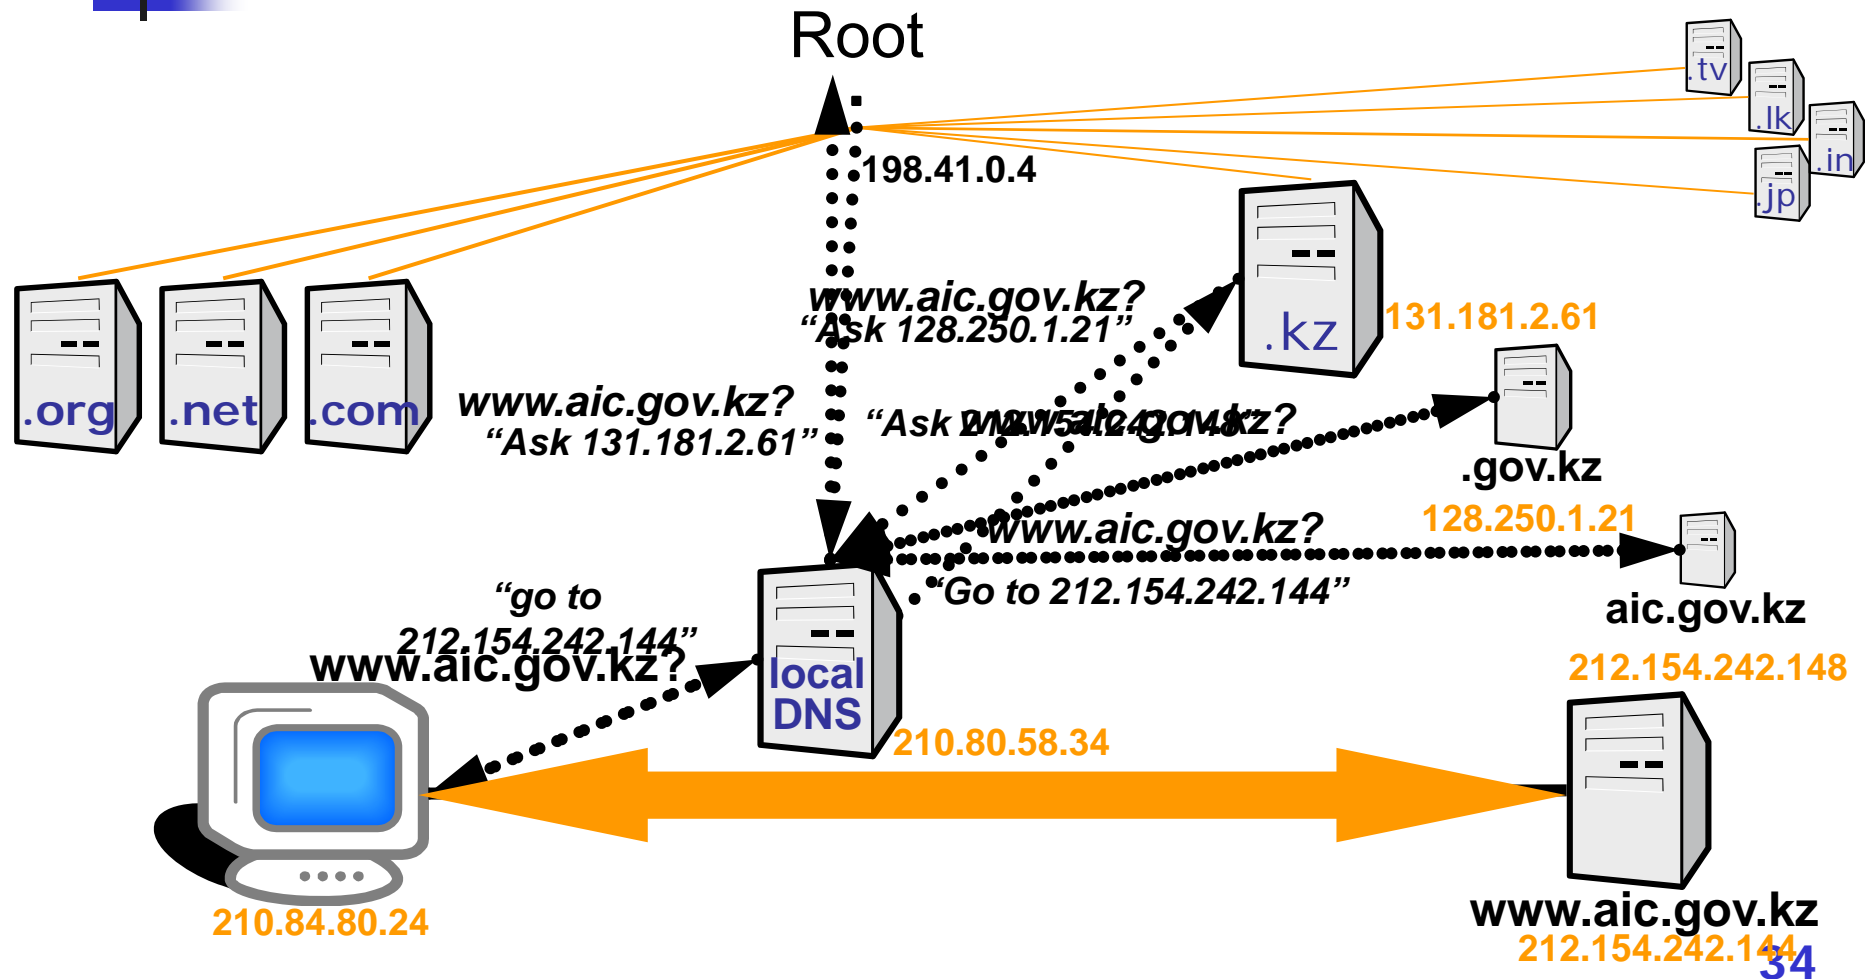


Mapping Domain Names to Addresses (6)

– example of iterative resolution




Internet addressing: an example of DNS




Mapping Domain Names to Addresses (7)

– caching mechanism to improve efficiency

- Caching at the name servers
- Caching at the hosts
- TLD servers typically cached in local name servers
 - Thus root name servers not often visited





Inverse Queries


- Mapping a particular resource to a domain name or domain names that have that resource
 - **Standard query**: mapping a domain name to a resource
 - **Inverse query**: mapping a **resource** to a **domain name**
- **Optional** part of DNS
- Generally **NOT** used because there is often no way to find the server that can resolve the query without searching the entire set of servers
- **NOT** an acceptable method of mapping host addresses to host names, use the **IN-ADDR.ARPA domain** instead



Pointer Queries

- Using IN-ADDR.ARPA domain for address to host mapping
 - An IP address in dotted-decimal format is included in the query
 - The correct domain name for the machine with the specific IP address
- **Note:** Since the IN-ADDR.ARPA special domain and the normal domain for a particular host or gateway will be in different zones, the possibility exists that that the data may be inconsistent





DNS Protocols



Related RFCs

- **RFC1034**, DOMAIN NAMES - CONCEPTS AND FACILITIES
- **RFC1035**, DOMAIN NAMES - IMPLEMENTATION AND SPECIFICATION
- Over 137 DNS related RFCs
 - See webpage at <http://www.dns.net/dnsrd/rfc/>



DNS Message Header Structure(1)

- **ID**: 16-bit field used to correlate queries and responses.
- **QR**: 1-bit field that identifies the message as a query (0) or response (1).
- **OPCODE**: 4-bit field that describes the type of query:
 - 0: Standard query (name to address). 1: Inverse query (address to name). 2: Server status request.
- **AA**: Authoritative Answer. 1-bit field. When set to 1, identifies this response is made by an authoritative name server.
- **TC**: Truncation. 1-bit field. When set to 1, indicates the message has been truncated due to length greater than that permitted.
- **RD**: Recursion Desired. 1-bit field. Set to 1 by the resolver to request recursive service by the name server.
- **RA**: Recursion Available. 1-bit field. Set to 1 by name server to indicate recursive query support is available.
- **Z**: 3-bit field. Reserved for future use. Must be set to 0.

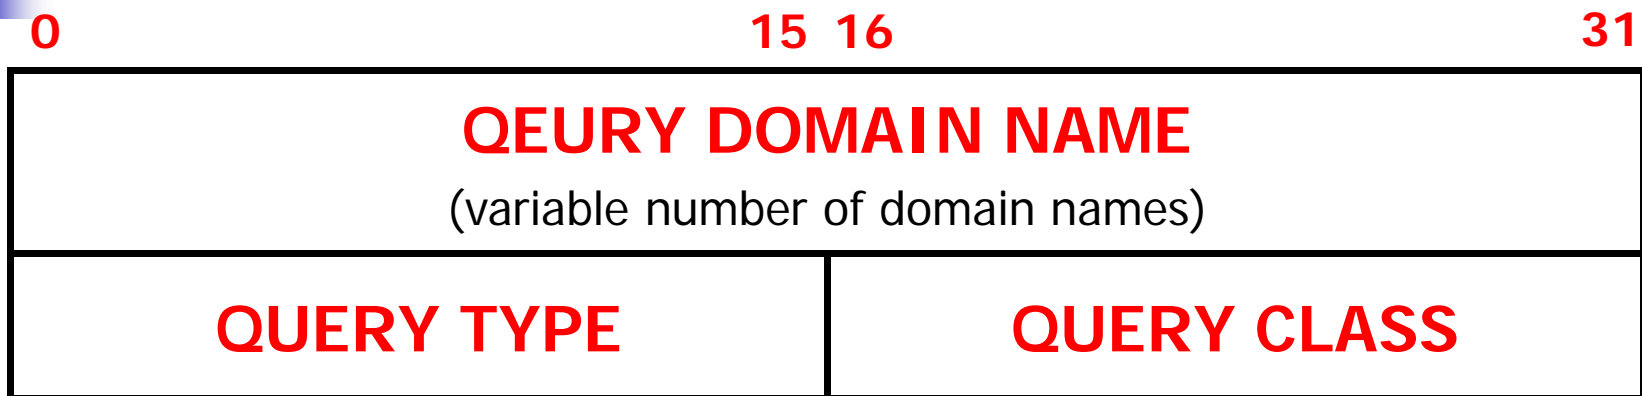


DNS Message Header Structure(2)

- **RCODE**: Response Code. 4-bit field that is set by the name server to identify the status of the query:
 - 0: No error condition. 1: Unable to interpret query due to format error. 2: Unable to process due to server failure. 3: Name in query does not exist. 4: Type of query not supported. 5: Query refused for policy reasons.
- **QDCOUNT(Question count)**: 16-bit field that defines the number of entries in the question section.
- **ANCOUNT(Answer count)**: 16-bit field that defines the number of resource records in the answer section.
- **NSCOUNT(Authority count)**: 16-bit field that defines the number of name server resource records in the authority section.
- **ARCOUNT(Additional count)**: 16-bit field that defines the number of resource records in the additional records section.



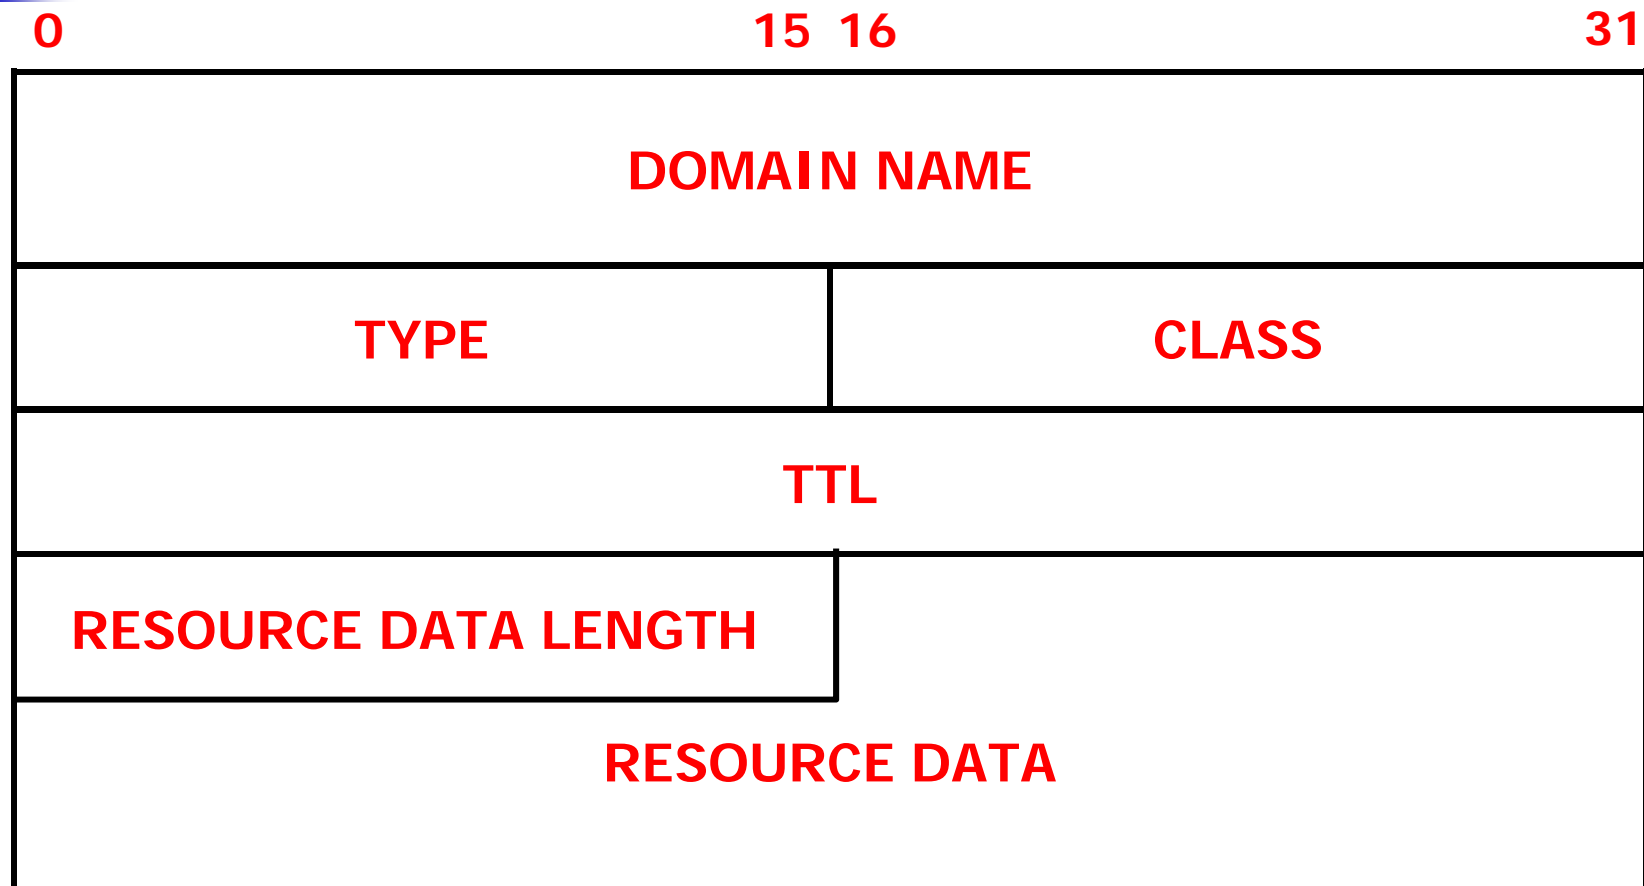
Question Section Format



- **QUERY TYPE**: 16-bit field used to specify the type of the query
 - A – Host address
 - MX – Mail exchanger for the domain
 - ...
- **QUERY CLASS**: 16-bit field used to specify the class of the query
 - IN – Internet system
 - ...



Resource Record Format (1)





Resource Record Format (2)

– type field

- SOA
 - Start Of Authority--identifies the domain or zone and sets a number of parameters
- NS
 - Maps a domain name to the name of a computer that is authoritative for the domain
- A
 - Maps the name of a system to its address. If a system (e.g., a router) has several addresses, then there will be a separate record for each.
- CNAME
 - Maps an alias name to the true, canonical name
- MX
 - Mail Exchanger. Identifies the systems that relay mail into the organization



Resource Record Format (3)

– type field

- TXT
 - Provides a way to add text comments to the database. For example, a txt record could map abc.com to the company's name, address, and telephone number
- WKS
 - Well Known Services. Can list the application services available at the host. Used sparingly, if at all
- HINFO
 - Host Information, such as computer type and model. Rarely used
- PTR
 - Maps an IP address to a system name. Used in address-to-name files.



Example of Resource Records(1)

- Type=A

- Name= Domain name , Value= IP Address

ns.bupt.edu.cn 86400 IN A 202.112.10.37

- Type=NS

- Name= Domain, eg. bupt.edu.cn
- value= Domain name of Authoritative Name Server

bupt.edu.cn 86400 IN NS ns.bupt.edu.cn



Example of Resource Records (2)

- Type=CNAME


- Name= canonical name
- Value= the name

dns.bupt.edu.cn 86400 IN CNAME ns.bupt.edu.cn

- Type=MX

- Name= Domain name
- Value= canonical name of mail server

bupt.edu.cn 86400 IN MX mail.bupt.edu.cn



DNS Tools



DNS Tools: nslookup

- Function: query Internet name servers [interactively](#)
- Examples:

```
[shiyang@localhost]$ nslookup mail.263.net
Server:                202.106.0.20
Address:               202.106.0.20#53

Non-authoritative answer:
Name:   mail.263.net
Address: 211.150.96.52
Name:   mail.263.net
Address: 211.150.96.51

[shiyang@localhost]$ nslookup 211.150.96.52
Server:                202.106.0.20
Address:               202.106.0.20#53
Non-authoritative answer:
52.96.150.211.in-addr.arpa      name = mail.263.net.
```



DNS Tools: dig

- Function: a flexible tool for interrogating DNS name servers
- Performs DNS lookups and displays the answers that are returned from the name server(s)

```
[shiyang@localhost]$ dig bupt.edu.cn
```

```
[shiyang@localhost]$ dig +norecurse bupt.edu.cn
```

```
[shiyang@localhost]$ dig +trace bupt.edu.cn
```

```
[shiyang@localhost]$ dig -x 211.68.71.130
```



Other Issues

- Dynamic DNS (DDNS)
 - RFC2136
- IDNS
 - Internationalized Domain Names (IDN)
- DNSSEC and other security issues
 - Security concern on DNS information exchange
- DNSv6
 - New frontier for next generation Internet
- ENUM
 - Convergence with telephony?
- More advanced topics
 - Digital Object Identifier, <http://www.doi.org/faq.html>
 - Handle System, <http://www.handle.net/>



Summary of Important Terms in DNS

- Domain / domain name
- Domain namespace
- Resource Record
- Name server
- Resolver
- Zone
- Query / response
- Standard query / inverse query / pointer query
- recursive resolution / iterative resolution
- Primary server/ secondary server / caching server



Some procedures should be understood

- How does DNS work together with the user programs (e.g. TELNET, FTP, HTTP) ?
- How is the recursive resolution and iterative resolution?
- What are the mechanisms in DNS that are possible to improve the querying efficiency?



Example of DNS Response(1) (RFC1034)

The query would look like:

Header	OPCODE=QUERY	
Question	QNAME=SRI-NIC.ARPA., QCLASS=IN, QTYPE=A	
Answer	<empty>	
Authority	<empty>	
Additional	<empty>	

The response from C.ISI.EDU would be:

Header	OPCODE=QUERY, RESPONSE, AA	
Question	QNAME=SRI-NIC.ARPA., QCLASS=IN, QTYPE=A	
Answer	SRI-NIC.ARPA. 86400 IN A 26.0.0.73 86400 IN A 10.0.0.51	
Authority	<empty>	
Additional	<empty>	

6.2.3. QNAME=SRI-NIC.ARPA, QTYPE=MX

This type of query might be result from a mailer trying to look up routing information for the mail destination HOSTMASTER@SRI-NIC.ARPA. The response from C.ISI.EDU would be:

Header	OPCODE=SQUERY, RESPONSE, AA			
Question	QNAME=SRI-NIC.ARPA., QCLASS=IN, QTYPE=MX			
Answer	SRI-NIC.ARPA. 86400 IN MX 0 SRI-NIC.ARPA.			
Authority	<empty>			
Additional	SRI-NIC.ARPA. 86400 IN A 26.0.0.73			
	A 10.0.0.51			



Example of DNS Response(3)

6.2.6. QNAME=BRL.MIL, QTYPE=A

If this query is sent to C.ISI.EDU, the reply would be:

Header	+-----+ OPCODE=SQUERY, RESPONSE +-----+			
Question	+-----+ QNAME=BRL.MIL, QCLASS=IN, QTYPE=A +-----+			
Answer	+-----+ <empty> +-----+			
Authority	+-----+ MIL. 86400 IN NS SRI-NIC.ARPA. +-----+			
	86400 NS A. ISI.EDU. +-----+			
Additional	+-----+ A. ISI.EDU. A 26.3.0.103 +-----+			
	SRI-NIC.ARPA. A 26.0.0.73 +-----+			
	A 10.0.0.51 +-----+			